

The Semantic Layer Is Not the Destination. It Is What Makes "Just Ask the Data" Trustworthy.

An implementation guide, with the sample metadata each phase actually produces.

Ask a database a plain question — *"how many customers did we have last month?"* — and it can hand you two different answers from the same rows. One counts customers; the other counts orders, so anyone who bought twice is counted twice. Both queries run, both look reasonable, and the wrong one is always the larger — so it is the one that reaches the slide. Swap "customers" for any figure your business actually reports and the trap is the same.

The same bug produced the two numbers this guide is about. Same warehouse, same afternoon, same question:

"On-time delivery for India warehouses last month."

88.5% and **87.3%**.

Same DuckDB file. Same SQL dialect. Same 132 orders. One of them is wrong, and in the review nobody could say which — which is the actual problem, and it is not a problem prompt engineering can reach.

Here is where 88.5% comes from. The `deliveries_raw` table has one row per delivery *leg*. An order that ships in two legs — partial shipment on Tuesday, remainder on Friday — has two rows. Join orders to legs and count, and that order votes twice. Worse: if the first leg arrived before the promised date, the order looks on time twice over, while the second leg that actually blew the promise is one row among 61.

61 legs. 55 orders.

- Order grain: **48 / 55 = 87.27%**
- Delivery grain: **54 / 61 = 88.52%**

The gap is 1.25 points. It is small enough to look like rounding, and it always errs upward. A wrong number that flatters you is the one that survives to the board deck.

The model did not hallucinate. It wrote syntactically valid, semantically reasonable SQL against a table whose grain nobody had written down anywhere it could read. **The failure was in the metadata, so the fix has to be in the metadata.**

The architectural rule

Everything below follows from one decision:

The model never writes SQL and never does arithmetic. It reads a governed metadata slice and emits a validated JSON object — which metric, which filters, which time window. A deterministic compiler turns that object into SQL.


This is the inversion that matters. Text-to-SQL asks a language model to be right about grain, joins, exclusions and fiscal calendars on every single call. This architecture asks it to be right about *one* much easier thing: which of three approved metrics you meant. Everything that must be exactly right is compiled from metadata, by code, the same way every time.

Two users asking the same question in different words get the same number, because the number was never in the model's hands.

The pipeline is six stages:

retriever → resolver → validator → compiler → executor → provenance

Only the **resolver** calls a model. Only the **compiler** writes SQL. And validation happens *before* compilation, so a refused question never becomes SQL at all.



Full resolution: [docs/diagrams/nl_dataflow.svg](#) — scaled to fit a page here, so open it directly to read the edge labels.

Each arrow is a stage reading a specific artifact for a specific reason. Two rows repay a second look: **the validator and the compiler never open a file**. Both receive a **SemanticModel** that the loader parsed, which means the compiler has no access to the glossary at all. Vocabulary is resolved *before* the gates, so which synonym a user typed can never reach the arithmetic.

Eight phases, nine artifacts


Each phase produces a real file and closes one specific way natural-language querying fails. That pairing is the whole method — a phase that doesn't close a named failure is a phase you can skip.

Phase	Artifact	The NL failure it closes
P0	00_kpi_contract.md	Nobody agrees what "on-time" means, so no answer can be wrong
P1	01_technical_metadata.json	Column names harvested from SAP DDIC, with no business meaning
P2	02_standardized_metadata.json	Types and keys inconsistent across source systems
P3	03_glossary.yml + 03_column_bindings.yml	"OTD", "delivery reliability", "promise adherence" reach no column
P4	04_process_model.yml	The model can answer <i>what</i> happened, never <i>why</i>
P5	05_semantic_model.yml	The model invents its own arithmetic — this is where 88.5% dies
P6	06_dq_rules.yml	A number computed from incomplete inputs is returned with full confidence
P7	07_catalog_asset.json	The answer arrives with no provenance, so nobody can defend it

Nine files, eight phases: P3 emits two, because a glossary term and its physical binding are different governance objects with different owners.

Where the metadata comes from, and where it goes

The question that lands next — usually from whoever is being asked to fund this — is "can't we just harvest all that from SAP?"



Full resolution: docs/diagrams/metadata_sources.svg — every box carries the collection method and the accountable owner, which is worth opening for.

Read it left to right: **source** → **how it was collected** → **the artifact** → **when the pipeline consumes it**. The colours are the answer to the funding question.

One of the nine artifacts is machine-harvestable. `01_technical_metadata.json` comes out of the SAP data dictionary via SE11 / SE16 — repeatable, cheap, automatable. The other eight split two ways:

Provenance	Count	What it means	Artifacts
HARVESTED	1	a repeatable query against a running system	P1
DERIVED	4	computed from earlier artifacts plus the data	P2, P3 bindings, P5, P7
AUTHORED	4	a named human decided something no system knows	P0, P3 glossary, P4, P6

The derived four are cheap to keep current, because their inputs are themselves versioned artifacts: P2 from P1 plus the raw rows plus staging SQL; the column bindings from the glossary checked against real warehouse columns at build time; P5 from P2's grain, P3's vocabulary and P4's phases; P7 from P5 + P6 plus governance fields.

The authored four are where the calendar time goes. Somebody had to decide that "on-time" means the *current* promised date and not the original (P0), that "OTIF" and "promise adherence" are sanctioned synonyms for the same metric (P3), that a standard QM usage decision takes one day (P4), and that 2% missing delivery dates is **degrading** rather than **blocking** (P6). No amount of crawling produces any of those, and none of them are discoverable from the data — they are judgements, made by someone accountable for them.

Which is the whole point, stated as a planning fact: **a harvester gets you tables, fields, domains and types. It cannot get you grain, exclusions, an approved definition or a standard duration** — and those are precisely the fields that decide whether the answer is right rather than merely well-formed. The 88.5% error lives in a field no crawler emits. This is why "point a catalog at the warehouse and turn on NL query" produces a demo and not a trustworthy system.

And then how it gets used

The right-hand band is the other half of the question, and it separates three things that get conflated:

- **P0 is never read at runtime.** It is the acceptance oracle — the thing the answer is judged *against*, not an input to producing it. If your KPI contract is being parsed by software, it has stopped being a contract.
- **P1 and P2 are consumed once, at build time,** by `sql/02_staging_model.sql` and the warehouse build. So the SAP data dictionary is not a live dependency of every question — which is both the truth and considerably less alarming than the alternative reading.
- **P3–P7 are read on every question,** by retriever, validator, compiler and provenance respectively.

That last line is the one that matters for cost and for failure modes: six artifacts sit in the hot path of every natural-language question, and if any of them is wrong, every answer is wrong the same way — silently, consistently, and with a provenance block attached that makes it look governed.

Both figures are generated from declared graphs, not drawn. Every arrow in the sourcing map carries a **probe** string that must occur in the artifact it points at, and the two diagrams are cross-checked against each other, so neither can claim a lineage the files don't support.

The synonym layer (P3)

Vocabulary is where NL querying quietly fails, and it is the cheapest phase to build. From the real file, with two catalog fields dropped:

```
- term: On-Time Delivery %
  definition: >
    Share of eligible customer orders delivered on or before the promised
    delivery date. Measured at order grain (one row per order_id). An order
    counts as delivered only when its final delivery leg arrives. Excludes
    cancelled orders and orders not yet due.
  synonyms: [OTD, on time delivery, on-time delivery, delivery performance,
    delivery reliability, promise adherence, OTIF, on time in full]
  owner: VP Supply Chain
  steward: Supply Chain Analytics Lead
```

Note that the definition states the grain *in prose the model reads*. That sentence — "one row per order_id" — is doing real work.

The governed metric (P5)

This is the load-bearing artifact. Abridged from the real file:

```
- name: on_time_delivery_pct
  label: On-Time Delivery %
  entity: fact_order_delivery
  grain: order
  time_dimension: promised_delivery_date
  expression:
    type: ratio_percent
    numerator: sum(CASE WHEN is_on_time THEN 1 ELSE 0 END)
    denominator: sum(CASE WHEN is_eligible THEN 1 ELSE 0 END)
  exclusions:
    - key: cancelled_orders
      predicate: order_status <> 'CANC'
      rationale: A cancelled order was never due; including it inflates the denominator.
    - key: not_yet_due
      predicate: promised_delivery_date <= DATE '2026-08-07'
      rationale: >
        An order promised after the as-of date is not late, it is pending.
  business_rules:
    - Measured against the CURRENT promised date, not the original.
    - Order-grain only. Delivery-grain aggregation double-counts split shipments.
  owner: VP Supply Chain
  approval_state: approved
  lineage: [SAP SD, SAP EWM, SAP TM, fact_order_delivery]
```

Four things to notice, because each one is a defect that would otherwise ship:

1. **grain: order** is declared, not inferred. This is the field that kills 88.5%.
2. **Exclusions carry a rationale**. Six months on, an exclusion without a stated reason gets deleted by someone who assumes it was a mistake.
3. **approval_state: approved** is read by the retriever, so an unapproved metric is never shown to the model. It cannot be asked for because it was never offered.
4. **Ratios are declared as numerator and denominator**, never as a pre-averaged rate. The compiler sums numerators and sums denominators, then divides. Averaging percentages across groups of unequal size gives a different, wrong number — and it looks fine.

The data-quality contract (P6)

Rules bind to specific metrics, so a defect degrades only what it actually touches:

```
- rule_id: dq_completeness_delivery_date
  dimension: completeness
  threshold: 2.0
  comparison: lte
  severity: degrading
  bound_to: [on_time_delivery_pct]
```

A trust badge — **TRUSTED**, **DEGRADED**, **BLOCKED** — ships with every answer. **BLOCKED** *replaces* the number rather than annotating it, because a caveated wrong number gets pasted into a slide with the caveat trimmed off.

The seven gates

Between the model's JSON and any SQL, seven checks run. All seven run every time, and a gate that cannot be evaluated **fails closed**:

Gate	Refuses
metric_approved	a metric that isn't approved, or doesn't exist
dimensions_declared	grouping by something the model invented
filters_bound	region = 'MARS' — a value outside the declared allowed set
grain_matches	the 88.5% query. It never becomes SQL.
no_fanout	a dimension only reachable through a one-to-many join
time_window_bounded	an unbounded or inverted date range
exclusions_applied	a metric whose certified exclusions can't be applied

A refusal names the gate and the reason. This is the actual output when the 88.5% query is put through the pipeline:

```
REFUSED on_time_delivery_pct: the intent did not pass validation.

gates failed grain_matches
grain_matches      Metric 'on_time_delivery_pct' is certified at grain 'order'; the
                    intent asks for grain 'delivery'. Aggregating at the wrong grain
                    changes the number.

No number is shown. A value that failed a gate is not a caveated answer,
it is an unverified one.
```

That is more useful than a number. "I can't answer that, and here is which rule stopped me" is a system you can put in front of a CFO. One that always produces a number is not.

Note **filters_bound** in particular. Without a declared allowed-value set, a filter on a misspelled region returns zero rows — and zero rows gets presented as an answer.

The question BI tools can't take

Descriptive questions ("what was it?") only need a glossary, a grain and a metric. The question people actually ask is diagnostic:

"Why did on-time delivery drop for India warehouses last month?"

That requires per-phase timestamps, per-phase planned baselines, and an attribution rule **defined once, in metadata**. This is why the manufacturing process is in scope: the answer to "why" lives upstream in production, not in the delivery table.

The process model declares twelve phases across SAP SD, PP, QM, MM/EWM, LE and TM — demand capture, MRP, production order release, material staging, production execution, quality inspection, goods receipt, picking, dispatch, transportation, delivery confirmation, billing. Nine of them are attributable spans. The other three aren't, and declaring them would let a filter pass every gate and return zero rows dressed as an answer.

For each order and phase: $\text{variance_days} = \text{actual_duration} - \text{standard_duration}$. The attribution is one **argmax** over the nine spans, ties broken to the earliest phase, evaluated only for orders that actually missed.

The real output, with the answer block's metadata lines elided:

ANSWER	On-Time Delivery % is 87.3% -- 48 of 55 eligible orders delivered on or before the promised date		
	... (metric, arithmetic, grain, window, filters, exclusions, lineage, trust, gates)		
BREAKDOWN			
	Quality inspection	4	
	Transportation	2	
	Production execution	1	
	unattributed	0	
	Largest contributor: Quality inspection (4 of 7 late orders).		
COMPARISON			
	prior window	June 2026 (2026-06-01 to 2026-06-30) on promised_delivery_date	
	prior value	96.2% (25 / 26)	
	change	down 8.9pp (96.2% -> 87.3%)	

unattributed 0 is printed even though it is zero, on purpose: an attribution that silently drops the orders it could not explain is an attribution you cannot audit.

On those 4 orders, the QM usage decision ran an average of **4.75 days over its 1.0-day standard, against 0.0 days of available slack** — the promise date left no room to absorb it. Plant 1010.

And the part that makes the case: **transportation ran under standard on 5 of the 7 late orders**. Logistics was the obvious suspect and logistics was faster than planned. Without per-phase baselines you get a plausible, confident, wrong story — and you go optimize the wrong department.

What this does *not* do

Worth stating plainly, because governance metadata invites overreading.

Catalog metadata supports **discovery**. It is not policy enforcement, consent management, access control, or retention automation. An **owner** field records who is accountable; it does not stop anyone from querying anything. Those are separate systems, and a semantic layer that implies otherwise is worse than one that stays silent.

Also honestly scoped: one fact table with declared joins, not a general warehouse; a frozen as-of date so the numbers are reproducible; a single-node engine. And the live resolver path is documented but **not** asserted as tested here — it needs API credentials this environment doesn't have, and the verification script reports that criterion as **SKIP** rather than quietly counting it as a pass.

Reproducing it

Everything above is measured, not illustrative. Five commands, no API key, no network:

```
pip install -r requirements.txt
python src/build_warehouse.py
python src/ask.py --offline Q1      # 87.3%, 48/55, TRUSTED, 7/7 gates
python src/ask.py --offline Q2      # the diagnostic breakdown above
python verify_acceptance.py         # measures every acceptance criterion
```

401 tests hold it in place, and the guide's prose is under test alongside the code — every figure in it is asserted against the artifacts or the warehouse, because a reproduction guide whose numbers have drifted spends the reader's trust before it spends their time.

If you would rather not clone anything: the companion **Medium post** carries the whole thing as a single self-contained file — embedded data, inline metadata, all seven gates, the compiler and the diagnostic — that runs on `pip install duckdb pyyaml` and reproduces every number above. It is the same file this repository tests, byte for byte.

The one-line version

Natural-language querying does not fail because models can't write SQL. It fails because the metadata that makes a question answerable *correctly* — grain, cardinality, exclusions, approved definitions, process baselines — was never written down where anything could read it.

Build that, and the model's job shrinks to something it is genuinely good at.

What is your team's answer to "which of these two numbers is right"? If it's "ask Priya," you have a semantic layer — it just isn't written down.